

# **A REPORT ON**

## **KAFKA Deployment And Creating Publisher/Consumer**

### **For ReST Based Services**



**International Institute Of Information  
Technology, Bangalore  
Computer Science and Engineering**

#### **Guidance by**

Prof. Samar Shailendra

#### **Submitted by**

Shubham Gupta (IMT2016118)

Suparn S Lele (MT2018122)

# **INDEX**

## **1. Abstract**

## **2. Introduction**

2.1 Why Dev-Ops

2.2 About Application

## **3. SDLC Phase**

3.1 Scope

3.2 Architecture

3.2.1 Jenkins in Docker Configuration

3.3 SCM

3.4 Build

3.5 Test

3.6 Deploy

3.7 Monitor

## **4. Results**

4.1 Scalability

## **5. Future Work**

## **6. Conclusion**

## **7. References**

## 1.0 Kafka Introduction:

**Apache Kafka** is an [open-source stream-processing](#) software platform developed by [LinkedIn](#) and donated to the [Apache Software Foundation](#), written in [Scala](#) and [Java](#). The project aims to provide a unified, high-throughput, low-latency platform for handling real-time data feeds.

Kafka uses a binary [TCP](#) design that is optimized for efficiency and relies on a "message set" abstraction that naturally groups messages together to reduce the overhead of the network roundtrip. This "leads to larger network packets, larger sequential disk operations, contiguous memory blocks [...]" which allows Kafka to turn a bursty stream of random message writes into linear writes."

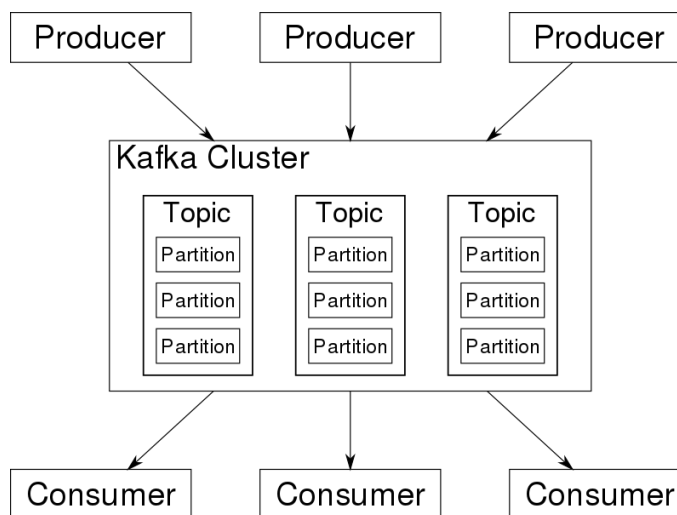


Fig 1. Kafka Architecture

## **2.0 Components :**

### **2.1 Zookeeper**

Apache Zookeeper is a distributed, open-source configuration,synchronization service along with naming registry for distributed applications. ZooKeeper stores a lot of shared information about [Kafka Consumers](#) and Kafka [Brokers](#).

### **2.2 Broker**

A Kafka cluster is made up of multiple Kafka Brokers. Each Kafka Broker has a unique ID (number). Kafka Brokers contain topic log partitions. Connecting to one broker bootstraps a client to the entire Kafka cluster. For failover, you want to start with at least three to five brokers. A Kafka cluster can have, 10, 100, or 1,000 brokers in a cluster if needed.

### **2.3 Producer**

Producers push data to brokers. When the new broker is started, all the producers search it and automatically sends a message to that new broker.

### **2.4 Consumer**

Consumers read data from brokers. Consumers subscribes to one or more topics and consume published messages by pulling data from the brokers.

### **2.5 Topic**

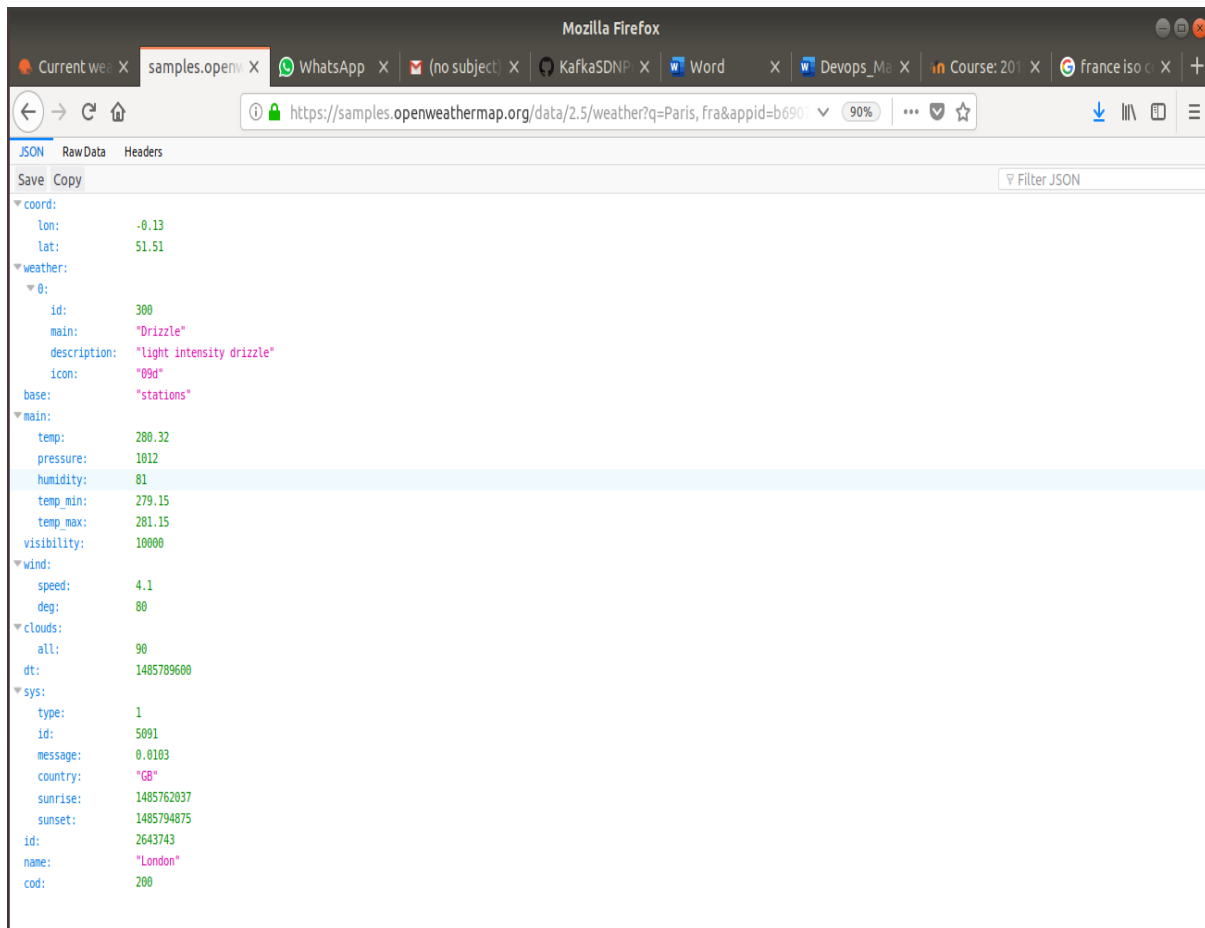
Kafka topic is a named stream of records. Kafka stores topics in logs. A topic log is broken up into partitions. Kafka spreads log's partitions across multiple servers or disks.

### 3.0 Data :

We are collecting weather data from <https://openweathermap.org/current#name>. Sample output from this website when we put London as a city is provided below.

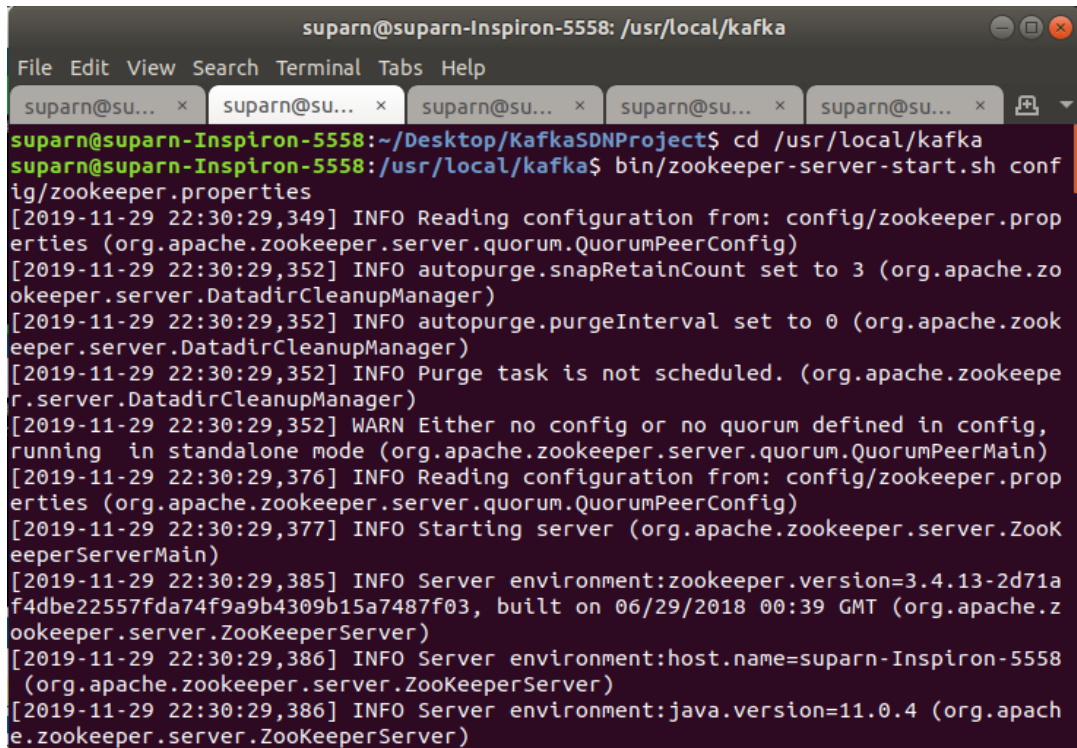
Link :-

<https://samples.openweathermap.org/data/2.5/weather?q=London,uk&appid=b6907d289e10d714a6e88b30761fae22>



## 4.0 Screenshots of Steps :

### Zoo-Keeper Server

A terminal window titled 'suparn@suparn-Inspiron-5558: /usr/local/kafka' showing the execution of 'bin/zookeeper-server-start.sh config/zookeeper.properties'. The output displays various log messages including configuration reading, autopurge settings, and server startup information. The terminal has a dark background with green and white text. The window title bar shows standard Linux window controls and the path '/usr/local/kafka'.

```
suparn@suparn-Inspiron-5558: /usr/local/kafka
File Edit View Search Terminal Tabs Help
suparn@su... x suparn@su... x suparn@su... x suparn@su... x suparn@su... x
suparn@suparn-Inspiron-5558:~/Desktop/KafkaSDNProject$ cd /usr/local/kafka
suparn@suparn-Inspiron-5558:/usr/local/kafka$ bin/zookeeper-server-start.sh config/zookeeper.properties
[2019-11-29 22:30:29,349] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2019-11-29 22:30:29,352] INFO autopurge.snapRetainCount set to 3 (org.apache.zookeeper.server.DataDirCleanupManager)
[2019-11-29 22:30:29,352] INFO autopurge.purgeInterval set to 0 (org.apache.zookeeper.server.DataDirCleanupManager)
[2019-11-29 22:30:29,352] INFO Purge task is not scheduled. (org.apache.zookeeper.server.DataDirCleanupManager)
[2019-11-29 22:30:29,352] WARN Either no config or no quorum defined in config, running in standalone mode (org.apache.zookeeper.server.quorum.QuorumPeerMain)
[2019-11-29 22:30:29,376] INFO Reading configuration from: config/zookeeper.properties (org.apache.zookeeper.server.quorum.QuorumPeerConfig)
[2019-11-29 22:30:29,377] INFO Starting server (org.apache.zookeeper.server.ZooKeeperServerMain)
[2019-11-29 22:30:29,385] INFO Server environment:zookeeper.version=3.4.13-2d71af4dbe22557fda74f9a9b4309b15a7487f03, built on 06/29/2018 00:39 GMT (org.apache.zookeeper.server.ZooKeeperServer)
[2019-11-29 22:30:29,386] INFO Server environment:host.name=suparn-Inspiron-5558 (org.apache.zookeeper.server.ZooKeeperServer)
[2019-11-29 22:30:29,386] INFO Server environment:java.version=11.0.4 (org.apache.zookeeper.server.ZooKeeperServer)
```

### Kafka Server

```
suparn@suparn-Inspiron-5558: /usr/local/kafka
File Edit View Search Terminal Tabs Help
suparn@su... x suparn@su... x suparn@su... x suparn@su... x suparn@su... x
suparn@suparn-Inspiron-5558:~/Desktop/KafkaSDNProject$ cd /usr/local/kafka
suparn@suparn-Inspiron-5558:/usr/local/kafka$ bin/kafka-server-start.sh config/s
erver.properties
[2019-11-29 22:30:38,381] INFO Registered kafka:type=kafka.Log4jController MBean
(kafka.utils.Log4jControllerRegistration$)
[2019-11-29 22:30:39,062] INFO starting (kafka.server.KafkaServer)
[2019-11-29 22:30:39,063] INFO Connecting to zookeeper on localhost:2181 (kafka.
server.KafkaServer)
[2019-11-29 22:30:39,087] INFO [ZooKeeperClient] Initializing a new session to l
ocalhost:2181. (kafka.zookeeper.ZooKeeperClient)
[2019-11-29 22:30:39,095] INFO Client environment:zookeeper.version=3.4.13-2d71a
f4dbe22557fda74f9a9b4309b15a7487f03, built on 06/29/2018 00:39 GMT (org.apache.z
ookeeper.ZooKeeper)
[2019-11-29 22:30:39,095] INFO Client environment:host.name=suparn-Inspiron-5558
(org.apache.zookeeper.ZooKeeper)
[2019-11-29 22:30:39,095] INFO Client environment:java.version=11.0.4 (org.apach
e.zookeeper.ZooKeeper)
[2019-11-29 22:30:39,096] INFO Client environment:java.vendor=Ubuntu (org.apache
.zookeeper.ZooKeeper)
[2019-11-29 22:30:39,096] INFO Client environment:java.home=/usr/lib/jvm/java-11
-openjdk-amd64 (org.apache.zookeeper.ZooKeeper)
[2019-11-29 22:30:39,096] INFO Client environment:java.class.path=/usr/local/kaf
ka/bin/../libs/activation-1.1.1.jar:/usr/local/kafka/bin/../libs/aopalliance-rep
ackaged-2.5.0-b42.jar:/usr/local/kafka/bin/../libs/argparse4j-0.7.0.jar:/usr/loc
```

